

South Dakota School of Mines and Technology

Advanced Topics in AI, Fall 2022

CSC 549 - M01

Exam 2

85%

1. Consider the following pseudocode for a faulty SARSA algorithm. Find all the mistakes in the algorithm. Explain why they are mistakes and correct them.

The image below shows the faults I have identified in the SARSA Algorithm.

```

procedure SARSA( number of episodes  $N \in \mathbb{N}$ 
                  discount factor  $\lambda \in (0, 1]$ 
                  learning rate  $\alpha_n = \frac{1}{\log(n+1)}$  ) ①
  Initialize matrices  $Q(s, a)$  and  $n(s, a)$  to 0,  $\forall s, a$ 
  for episode  $k \in 1, 2, 3, \dots, n$  do
     $t \leftarrow 1$ 
    Initialize  $s_1$ 
    Choose  $a_1$  from a uniform distribution over the actions ②
    while Episode  $k$  is not finished do
      Take action  $a_t$ : observe reward  $r_t$  and next state  $s_{t+1}$ 
      Choose  $a_{t+1}$  from  $s_{t+1}$  using  $\mu_t$ : an  $\epsilon$ -greedy policy with respect to  $Q$ 
      if The current state is terminal then  $\triangleright$  Compute target value
         $y_t = 0$  Terminal reward should be  $r$ .
      else
         $y_t = r_t + \max_a Q(s_{t+1}, a)$  This is off-policy, but Sarsa should be on-policy
      end if
       $n(s_t, a_t) \leftarrow n(s_t, a_t) + 1$ 
      Update Q function:
        ③  $Q(s_{t+1}, a_{t+1}) \leftarrow Q(s_t, a_t) - \alpha_{n(s_t, a_t)} (y_t - Q(s_t, a_t))$ 
       $t \leftarrow t + 1$ 
    end while
  end for
end procedure

```

- (i) The learning rate ($\alpha_n = \frac{1}{\log(n+1)}$) would be greater than 1 if the base of the log term is not **2**.
- (ii) Choosing a_1 from a uniform distribution over the actions would result in an off-policy action which is not what we want. A more appropriate choice would be to **choose a_1 from s_1 using the policy μ_t : an ϵ -greedy policy w.r.t. Q** .
- (iii) This algorithm contains all the characteristics of a SARSA (λ) algorithm; so I feel it is safe to assume that this algorithm was designed to replicate a SARSA (λ) algorithm. What is missing here is the backtracking update (**Eligibility Trace Update**). The corrections for (3) are shown below:

For all $s \in S, a \in A(s)$:

$$Q(s, a) \leftarrow Q(s, a) + \alpha y_t n(s, a)$$

$$n(s, a) \leftarrow \lambda n(s, a)$$

* here we assumed that $\lambda = 1$

2. Your friend found a variant of SARSA which is defined through a sequence of policies π_t (where $t \geq 1$), and consist of just changing (in the previous algorithm after corrections) the way the target is computed. The target becomes:

$$y_t = r_t + \lambda \sum_a \pi_t(a | s_{t+1}) Q(S_{t+1}, a)$$

Where $\pi_t(a | s)$ is the probability that a is selected in state s under policy π_t .

- a) What sequence of policies (π_t) should you choose so that the corresponding variant of SARSA is on-policy? The variant is called Expected SARSA.

- A. Expected SARSA would still be on-policy if the chosen action (a) by π_t resulted in the largest value of $\pi_t(a | s) Q(S_{t+1}, a)$ when compared to other actions in the state. Or in other words, a chosen should reflect the dominant $\pi_t(a | s) Q(S_{t+1}, a)$ in the term $\sum_a \pi_t(a | s_{t+1}) Q(S_{t+1}, a)$. This would only happen if the sequence of policies followed are determined by an ϵ -greedy policy w.r.t. Q where $\epsilon = 1$.

- b) Consider an off policy variant of SARSA corresponding to a stationary policy $\pi = \pi_t \forall t$. Under this algorithm, do the Q values converge? If so, what are the limiting Q values? Justify your answer.

- A. In the trivial case, if the stationary policy being followed doesn't reach the terminal state, the Q values will not converge as states which are looped over would have a Q value of $-\infty$.

The Q values would converge and would be dictated by the Q value of the terminal state. The Q values for each state would converge to the sum of the Q values of the next state plus the reward for taking the action that leads to that next state. For example, the Q value of the state just before the Terminal State in the policy would

have a limiting Q value of the Terminal State (0) plus the the reward for taking the action that led to the Terminal State (-1) which equals to -1.

My thought process behind this was based of value iteration from Chapter 4. When the episodes reach infinity, the update $(\alpha(y_t - Q(s_t, a_t)))$ should tend to 0. And from observing the grid world problem, the stable value of each state in the policy was somewhat of the form:

$$Q(s_t, a_t) = Q(s_{t+1}, a_t) + r_t$$

And since the terminal states Value function is fixed at 0, it dictates the values of all preceding states by this equation.