

# CSC 449 Advanced Topics in Artificial Intelligence

## Exam 2

Nov 30, 2022

Patrick McBride

90

### Problem 1:

Consider the following pseudo-code for a faulty SARSA algorithm:

```
procedure SARSA( number of episodes  $N \in \mathbb{N}$ 
                  discount factor  $\lambda \in (0, 1]$ 
                  learning rate  $\alpha_n = \frac{1}{\log(n+1)}$  )
  Initialize matrices  $Q(s, a)$  and  $n(s, a)$  to 0,  $\forall s, a$ 
  for episode  $k \in 1, 2, 3, \dots, n$  do
     $t \leftarrow 1$ 
    Initialize  $s_1$ 
    Choose  $a_1$  from a uniform distribution over the actions
    while Episode  $k$  is not finished do
      Take action  $a_t$ : observe reward  $r_t$  and next state  $s_{t+1}$ 
      Choose  $a_{t+1}$  from  $s_{t+1}$  using  $\mu_t$ : an  $\epsilon$ -greedy policy with respect to  $Q$ 
      if The current state is terminal then
         $y_t = 0$ 
      else
         $y_t = r_t + \max_a Q(s_{t+1}, a)$ 
      end if
       $n(s_t, a_t) \leftarrow n(s_t, a_t) + 1$ 
      Update Q function:
         $Q(s_{t+1}, a_{t+1}) \leftarrow Q(s_t, a_t) - \alpha_{n(s_t, a_t)} (y_t - Q(s_t, a_t))$ 
       $t \leftarrow t + 1$ 
    end while
  end for
end procedure
```

**Error 1** points to the learning rate  $\alpha_n = \frac{1}{\log(n+1)}$ .

**Error 2** points to the line "Choose  $a_1$  from a uniform distribution over the actions".

**Error 3** points to the Q function update line.

Annotations in yellow boxes:

- Next to  $y_t = 0$ : "should be r"
- Next to  $y_t = r_t + \max_a Q(s_{t+1}, a)$ : "no max, and missing lambda"

Find all of the mistakes in the algorithm. Explain why they are mistakes, and correct them.

- 1) This adaptive learning rate is ok so long as log is base 2 since alpha must be bounded by (0,1]
- 2) Basic SARSA( $\lambda$ ) is an on-policy TD control algorithm, so the action  $a_t$  must be chosen from  $s_t$  using a policy derived from  $Q(\epsilon - \text{greedy})$  (AKA  $\mu_t$ ).
- 3) After confirming that the intended implementation for this pseudo-code is SARSA( $\lambda$ ). The update to  $Q$  should include eligibility traces. (Since we were not given a  $\lambda$  we will assume it is 1). So the  $Q$  update should be:

$\forall s \in S, a \in A(s):$

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha_t n(s_t, a_t) [y_t - Q(s_t, a_t)]$$

$$n(s_t, a_t) \leftarrow \lambda n(s_t, a_t)$$

## Problem 2:

Your friend found a variant of SARSA which is defined through a sequence of policies  $\pi_t$  (where  $t \geq 1$ ), and consists of just changing (in the previous algorithm after corrections) the way the target is computed. The target becomes

$$y_t = r_t + \lambda \sum \pi_t(a|s_{t+1})Q(S_{t+1},a),$$

where  $\pi_t(a|s)$  is the probability that  $a$  is selected in state  $s$  under policy  $\pi_t$ .

**A. What sequence of policies ( $\pi_t$ ) should you choose so that the corresponding variant of SARSA is on-policy? This variant is called Expected SARSA.**

For a RL algorithm to be on-policy it must follow the policy that it is learning when choosing actions. This means that for Expected SARSA the action that is taken must be that action with the highest value in the current state. The action taken will be decided by a deterministic policy that is  $\epsilon$ -greedy with respect to the current Q function. In other words:  $\forall t$ ,  $\pi_t$  will be replaced by  $\pi_{t+1}$  which is an  $\epsilon$ -greedy policy with respect to the current Q function.

**B. Consider an off-policy variant of SARSA corresponding to a stationary policy  $\pi = \pi_t \forall t$ . Under this algorithm, do the Q values converge? If so, what are the limiting Q values? Justify your answer.**

If following  $\pi_t$  will result in eventually reaching the terminal state:

The Q values will converge. They will be limited by the Q value of the terminal state. This is because the Q value of the terminal state will never change as it is never left. Therefore all other states will have their Q values based off the Q value of the terminal state.

If following  $\pi_t$  will never result in reaching the terminal state:

The Q values will not converge since upon each update they will be driven further negative. As the number of iterations grows the Q values of the states that we visit will eventually be updated by  $-\infty$  each time.