

Gradient Descent

Michail Michailidis & Patrick Maiden


Outline

- Motivation
- Gradient Descent Algorithm
 - Issues & Alternatives
- Stochastic Gradient Descent
- Parallel Gradient Descent
- HOGWILD!

Motivation

- It is good for finding global minima/maxima if the function is convex
- It is good for finding local minima/maxima if the function is not convex
- It is used for optimizing many models in Machine learning:
 - **It is used in conjunction with:**
 - Neural Networks
 - Linear Regression
 - Logistic Regression
 - Back-propagation algorithm
 - Support Vector Machines


Function Example




Quickest ever review of multivariate calculus

- Derivative
- Partial Derivative
- Gradient Vector


Derivative



$$f(x) = x^2$$



$$f'(x) = df/dx = 2x$$



$$f''(x) = d^2 f/dx^2 = 2$$

- Slope of the tangent line





Figure 6.2

- Easy when a function is univariate


Partial Derivative – Multivariate Functions

For multivariate functions (e.g two variables) we need partial derivatives – one per dimension. Examples of multivariate functions:




$$f(x,y) = x^2 + y^2$$

Convex!




$$f(x,y) = -x^2 - y^2$$

Concave!



$$f(x,y) = \cos^2(x) + y^2$$




$$f(x,y) = \cos^2(x) + \cos^2(y)$$

Partial Derivative – Cont'd


To visualize the partial derivative for each of the dimensions x and y , we can imagine a plane that “cuts” our surface along the two dimensions and once again we get the slope of the tangent line.

The surface defined by $f(x,y) = 9 - x^2 - y^2$.




surface: $f(x,y) = 9 - x^2 - y^2$

Adding the plane $y = 1$.



plane: $y = 1$

The tangent line in the direction of x .




cut: $f(x,1) = 8 - x^2$

slope / derivative of cut: $f'(x) = -2x$

Partial Derivative – Cont'd 2

If we partially differentiate a function with respect to x, we pretend y is constant




$$f(x, y) = 9 - x^2 - y^2$$

$$f(x, y) = 9 - x^2 - c^2$$

$$f(x, y) = 9 - c^2 - y^2$$


$$f \downarrow x = \frac{\partial f}{\partial x} = -2x$$

$$f \downarrow y = \frac{\partial f}{\partial y} = -2y$$



Partial Derivative – Cont'd 3

The two tangent lines that pass through a **point**, define the tangent plane to that **point**



Gradient Vector

- Is the vector that has as coordinates the partial derivatives of the function:


$$f(x,y)=9-x^2-y^2 \qquad \frac{\partial f}{\partial x}=-2x \quad \frac{\partial f}{\partial y}=-2y$$

$$\nabla f = \frac{\partial f}{\partial x} i + \frac{\partial f}{\partial y} j = (\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}) = (-2x, -2y)$$

- **Note: Gradient Vector is not parallel to tangent surface**

Gradient Descent Algorithm & Walkthrough

- Idea
 - Start somewhere
 - Take steps based on the gradient vector of the current position till convergence
- Convergence :
 - happens when change between two steps $< \epsilon$



Gradient Descent Code (Python)


```
# From calculation, we expect that the local minimum occurs at x=9/4
```

```
x_old = 0
x_new = 6 # The algorithm starts at x=6
eps = 0.01 # step size
precision = 0.00001
```


```
def f_prime(x):
    return 4 * x**3 - 9 * x**2

while abs(x_new - x_old) > precision:
    x_old = x_new
    x_new = x_old - eps * f_prime(x_old)
print "Local minimum occurs at ", x_new
```

$$f'(x) = 4x^3 - 9x^2$$

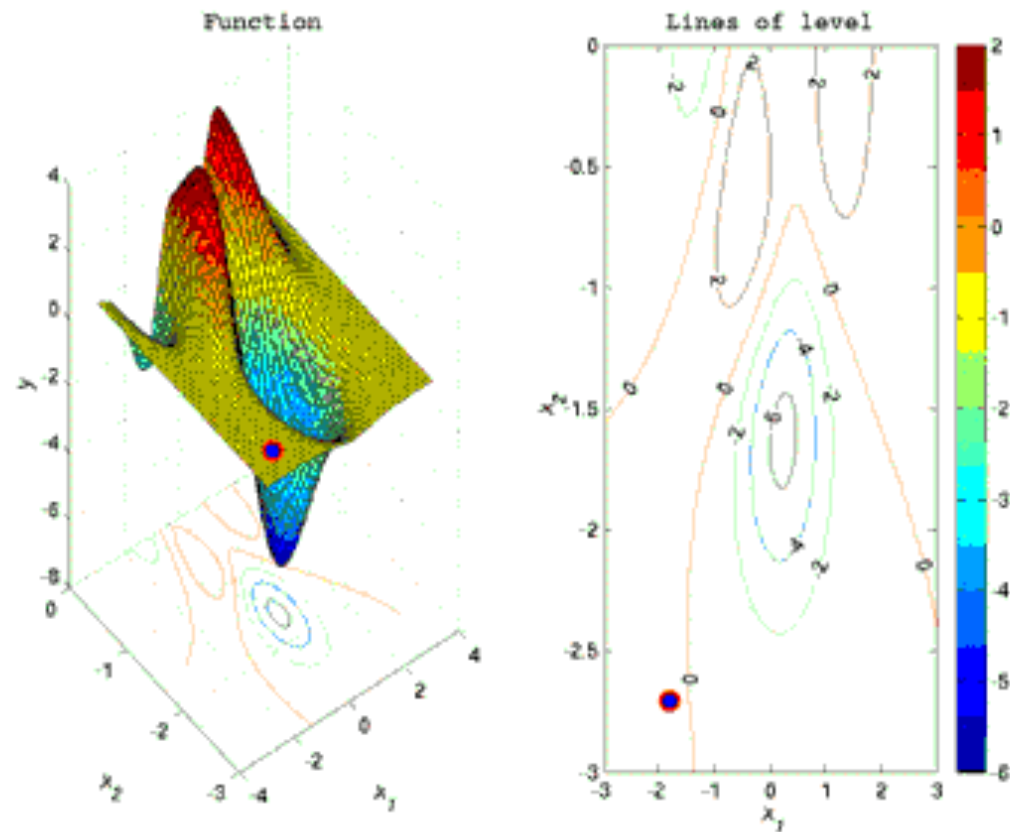


$$f(x) = x^4 - 3x^3 + 2$$



$$f'(x) = 4x^3 - 9x^2$$

Gradient Descent Algorithm & Walkthrough






Potential issues of gradient descent - Convexity




We need a **convex** function
→ so there is a global minimum:

$$f(x,y)=x^2 + y^2$$



Potential issues of gradient descent – Convexity (2)



Potential issues of gradient descent – Step Size

- As we saw before, one parameter needs to be set is the step size
- Bigger steps leads to faster convergence, right?

Alternative algorithms

- Newton's Method
 - Approximates a polynomial and jumps to the min of that function
 - Needs Hessian
- BFGS
 - More complicated algorithm
 - Commonly used in actual optimization packages

Stochastic Gradient Descent

- Motivation

- One way to think of gradient descent is as a minimization of a sum of functions:

- $w = w - \alpha \nabla L(w) = w - \alpha \sum_i \nabla L_i(w)$

- (L_i is the loss function evaluated on the i -th element of the dataset)

- On large datasets, it may be computationally expensive to iterate over the whole dataset, so pulling a subset of the data may perform better
- Additionally, sampling the data leads to “noise” that can avoid finding “shallow local minima.” This is good for optimizing non-convex functions. (Murphy)

Stochastic Gradient descent

- Online learning algorithm
- Instead of going through the entire dataset on each iteration, randomly sample and update the model

Initialize w and α

Until convergence do:

 Sample one example i from dataset //stochastic portion

$w = w - \alpha \nabla L_i(w)$

return w

Stochastic Gradient descent (2)

- Checking for convergence after each data example can be slow
- One can simulate stochasticity by reshuffling the dataset on each pass:

Initialize w and α

Until convergence do:

 shuffle dataset of n elements //simulating stochasticity

 For each example i in n :

$$w = w - \alpha \nabla L_i(w)$$

return w

- This is generally faster than the classic iterative approach (“noise”)
- However, you are still passing over the entire dataset each time
- An approach in the middle is to sample “batches”, subsets of the entire dataset
 - This can be parallelized!

Parallel Gradient descent

- Training data is chunked into batches and distributed

Initialize w and α

Loop until convergence:

 generate randomly sampled chunk of data m

 on each worker machine v :

$\nabla L \downarrow v(w) = \text{sum}(\nabla L \downarrow i(w))$ // compute gradient on batch

$w = w - \alpha * \text{sum}(\nabla L \downarrow v(w))$ //update global w model

return w

HOGWILD! (Niu, et al. 2011)

- Unclear why it is called this
- Idea:
 - In Parallel SGD, each batch needs to finish before starting next pass
 - In HOGWILD!, share the global model amongst all machines and update on-the-fly
 - No need to wait for all worker machines to finish before starting next epoch
 - Assumption: component-wise addition is atomic and does not require locking

HOGWILD! - Pseudocode

Initialize global model w

On each worker machine:

 loop until convergence:

 draw a sample e from complete dataset E

 get current global state w and compute $\nabla L_e(w)$

 for each component i in e :

$$w_i = w_i - \alpha b_v^T \nabla L_e(w) \quad // \quad b_v \text{ is } v^{\text{th}} \text{ std. basis}$$


component

 update global w


return w

Comparison

Parallel SGD



HOGWILD!



Comparison


type	data set	size (GB)	ρ	Δ	HOGWILD!			ROUND ROBIN		
					time (s)	train error	test error	time (s)	train error	test error
SVM	RCV1	0.9	0.44	1.0	9.5	0.297	0.339	61.8	0.297	0.339
MC	Netflix	1.5	2.5e-3	2.3e-3	301.0	0.754	0.928	2569.1	0.754	0.927
	KDD	3.9	3.0e-3	1.8e-3	877.5	19.5	22.6	7139.0	19.5	22.6
	Jumbo	30	2.6e-7	1.4e-7	9453.5	0.031	0.013	N/A	N/A	N/A
Cuts	DBLife	3e-3	8.6e-3	4.3e-3	230.0	10.6	N/A	413.5	10.5	N/A
	Abdomen	18	9.2e-4	9.2e-4	1181.4	3.99	N/A	7467.25	3.99	N/A

Figure 2: Comparison of wall clock time across of HOGWILD! and RR. Each algorithm is run for 20 epochs and parallelized over 10 cores.

- RR – Round Robin
 - Each machine updates x as it comes in. Wait for all before starting next pass
- AIG
 - Like Hogwild but does fine-grained locking of variables that are going to be used

Comparison (2)

SVM



Graph Cuts





Figure 3: Total CPU time versus number of threads for (a) RCV1, (b) Abdomen, and (c) DBLife.

Matrix Completion

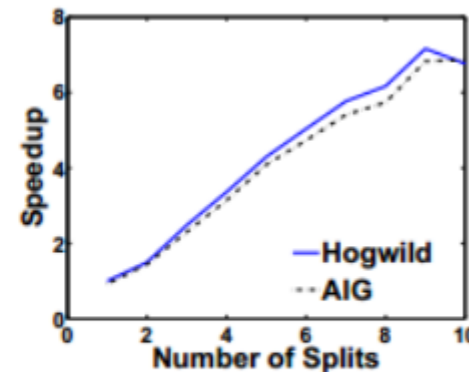

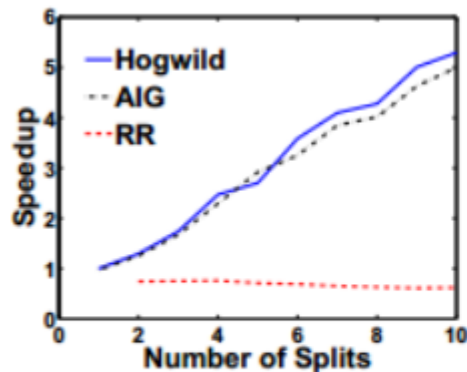


Figure 4: Total CPU time versus number of threads for the matrix completion problems (a) Netflix Prize, (b) KDD Cup 2011, and (c) the synthetic Jumbo experiment.

Moral of the story

- Having an idea of how gradient descent works informs your use of others' implementations
- There are very good implementations of the algorithm and other approaches to optimization in many languages
- Packages:
 - Python
 - [NumPy/SciPy](#)
 - Matlab
 - [Matlab Optimization toolbox](#)
 - [Pmtk3](#)
 - R
 - General-purpose optimization: [optim\(\)](#)
 - [R Optimization Infrastructure \(ROI\)](#)
 - TupleWare
 - Coming soon....

Resources

Partial Derivatives:

- <http://msemac.redwoods.edu/~darnold/math50c/matlab/pderiv/index.xhtml>
- http://mathinsight.org/nondifferentiable_discontinuous_partial_derivatives
- <http://www.sv.vt.edu/classes/ESM4714/methods/df2D.html>
- Gradients Vector Field Interactive Visualization: <http://dlippman.imathas.com/g1/Grapher.html> from https://www.khanacademy.org/math/calculus/partial_derivatives_topic/gradient/v/gradient-1
- <http://simmakers.com/wp-content/uploads/Soft/gradient.gif>

Gradient Descent:

- http://en.wikipedia.org/wiki/Gradient_descent
- http://www.youtube.com/watch?v=5u4G23_OoI (Stanford ML Lecture 2)
- http://en.wikipedia.org/wiki/Stochastic_gradient_descent
- Murphy, *Machine Learning, a Probabilistic Perspective*, 2012, MIT Press
- Hogwild paper: <http://pages.cs.wisc.edu/~brecht/papers/hogwildTR.pdf>